
Hydrodynamic simulations on a moving Voronoi mesh

Volker Springel

Heidelberg Institute for Theoretical Studies, Schloss-Wolfsbrunnenweg 35, 69118 Heidelberg, Germany, volker.springel@h-its.org

Zentrum für Astronomie der Universität Heidelberg, ARI, Mönchhofstr. 12-14, 69120 Heidelberg, Germany, volker.springel@uni-hd.de

Summary At the heart of any method for computational fluid dynamics lies the question of how the simulated fluid should be discretized. Traditionally, a fixed Eulerian mesh is often employed for this purpose, which in modern schemes may also be adaptively refined during a calculation. Particle-based methods on the other hand discretize the mass instead of the volume, yielding an approximately Lagrangian approach. It is also possible to achieve Lagrangian behavior in mesh-based methods if the mesh is allowed to move with the flow. However, such approaches have often been fraught with substantial problems related to the development of irregularity in the mesh topology. Here we describe a novel scheme that eliminates these weaknesses. It is based on a moving unstructured mesh defined by the Voronoi tessellation of a set of discrete points. The mesh is used to solve the hyperbolic conservation laws of ideal hydrodynamics with a finite volume approach, based on a second-order Godunov scheme with an exact Riemann solver. A particularly powerful feature of the approach is that the mesh-generating points can in principle be moved arbitrarily. If they are given the velocity of the local flow, a highly accurate Lagrangian formulation of continuum hydrodynamics is obtained that is free of mesh distortion problems, while it is at the same time fully Galilean-invariant, unlike ordinary Eulerian codes. We describe the formulation and implementation of our new Voronoi-based hydrodynamics, and we discuss a number of illustrative test problems that highlight its performance in practical applications.

1 Introduction

Numerical simulations have become an indispensable tool to study fluid dynamics, especially in astrophysics where direct experiments are often impossible. However, it is not always clear whether the employed simulation algorithms are sufficiently accurate in real practical applications, and to which

extent numerical deficits may affect the final results. It therefore remains an important task to continue to critically test the numerical methods that are in use, and to develop new approaches with the goal to reach better accuracy at comparable or even lower computational cost.

In astrophysics, a variety of fundamentally quite different numerical methods for hydrodynamical simulations are in use, the most prominent ones are smoothed particle hydrodynamics (SPH; Lucy, 1977; Gingold & Monaghan, 1977; Monaghan, 1992; Springel, 2010b) and Eulerian mesh-based hydrodynamics (e.g. Toro, 1997; LeVeque, 2002; Stone et al., 2008) with (optional) adaptive mesh refinement (AMR). A particular challenge in astronomy is the need to calculate self-gravitating flows, which often tend to cluster strongly under gravity, producing a huge dynamic range in density and length scales that can only be treated efficiently with spatially adaptive resolution. An important reason for the popularity of SPH lies in the fact that such an adaptivity is automatically built into this method, whereas achieving it in adaptive mesh refinement codes requires substantial effort.

It has become clear over recent years that both SPH and AMR suffer from fundamental problems that make them inaccurate in certain regimes. Indeed, these methods sometimes yield conflicting results even for basic calculations that only consider non-radiative hydrodynamics (e.g. Frenk et al., 1999; Agertz et al., 2007; Tasker et al., 2008; Mitchell et al., 2009). SPH codes have comparatively poor shock resolution, offer only low-order accuracy for the treatment of contact discontinuities, and suffer from subsonic velocity noise (Abel, 2011). Worse, they appear to suppress fluid instabilities under certain conditions (Ageritz et al., 2007), as a result of a spurious surface tension and inaccurate gradient estimates across density jumps. On the other hand, Eulerian codes are not free of fundamental problems either. They do not produce Galilean-invariant results, which can make their accuracy sensitive to the presence of bulk velocities (e.g. Wadsley et al., 2008; Tasker et al., 2008). Another concern lies in the mixing inherent in multi-dimensional Eulerian hydrodynamics. This provides for an implicit source of entropy, with sometimes unclear consequences (e.g. Wadsley et al., 2008).

There is hence substantial motivation to search for new hydrodynamical methods that improve on these weaknesses of the SPH and AMR techniques. In particular, we would like to retain the accuracy of mesh-based hydrodynamical methods (for which decades of experience have been accumulated in computational fluid dynamics), while at the same time we would like to outfit them with the Galilean-invariance and geometric flexibility that is characteristic of SPH. The principal idea for achieving such a synthesis is to allow the mesh to move with the flow itself. This is an obvious and old idea (Braun & Sambridge, 1995; Gnedin, 1995; Whitehurst, 1995; Mavriplis, 1997; Xu, 1997; Hassan et al., 1998; Pen, 1998; Trac & Pen, 2004), but one fraught with many practical difficulties that have so far prevented widespread use of any of the few past attempts to introduce moving-mesh methods in astrophysics and cosmology. For example, Gnedin (1995) and Pen (1998) presented moving-




Fig. 1. Example of a Voronoi and Delaunay tessellation in 2D, with periodic boundary conditions. The red circles show the generating points of the Voronoi tessellation, which is drawn with solid lines. Its topological dual, the Delaunay triangulation, is overlaid with thin dashed lines.

mesh hydrodynamic algorithms which relied on the continuous deformation of a Cartesian grid. However, the need to limit the maximum allowed grid distortions severely impacts the flexibility of these codes for situations in which the mesh becomes heavily distorted, and special measures are required to let the codes evolve cosmological density fields into a highly clustered state. In general, mesh tangling (manifested in ‘bow-tie’ cells and hourglass like mesh motions) is the traditional problem of such attempts to simulate multi-dimensional hydrodynamics in a Lagrangian fashion.

In this contribution, we describe a new formulation of continuum hydrodynamics based on an unstructured mesh. The mesh is defined as the Voronoi tessellation of a set of discrete mesh-generating points, which are in principle allowed to move freely. For the given set of points, the Voronoi tessellation of space consists of non-overlapping cells around each of the sites such that each cell contains the region of space closer to it than to any of the other sites. Closely related to the Voronoi tessellation is the Delaunay tessellation,

the topological dual of the Voronoi diagram. Both constructions have already been widely used for natural neighbor interpolation and geometric analysis of cosmic structures (e.g. van de Weygaert, 1994; Sambridge et al., 1995; Schaap & van de Weygaert, 2000; Pelupessy et al., 2003; van de Weygaert & Schaap, 2009). In 2D, the Delaunay tessellation for a given set of points is a triangulation of the plane, where the points serve as vertices of the triangles. The defining property of the Delaunay triangulation is that each circumcircle around one of the triangles of the tessellation is not allowed to contain any of the other mesh-generating points in its interior. This empty circumcircle property distinguishes the Delaunay triangulation from the many other triangulations of the plane that are possible for the point set, and in fact uniquely determines the triangulation for points in general position. Similarly, in three dimensions, the Delaunay tessellation is formed by tetrahedra that are not allowed to contain any of the points inside their circumspheres.

As an example, Figure 1 shows the Delaunay and Voronoi tessellations for a small set of points in 2D, enclosed in a box with imposed periodic boundary conditions. The midpoints of the circumcircles around each Delaunay triangle form the vertices of the Voronoi cells, and for each line in the Delaunay diagram, there is an orthogonal face in the Voronoi tessellation.

The Voronoi cells can be used as control volumes for a finite-volume formulation of hydrodynamics, using the same principal ideas for reconstruction, evolution and averaging (REA) steps that are commonly employed in many Eulerian techniques. However, as we will see it is possible to consistently include the mesh motion in the formulation of the numerical steps, allowing the REA-scheme to become Galilean-invariant. Even more importantly, due to the mathematical properties of the Voronoi tessellation, the mesh continuously deforms and changes its topology as a result of the point motion, without ever leading to the dreaded mesh-tangling effects that are the curse of traditional moving mesh methods. We note that the approach we describe here is quite different from attempts to formulate fluid particle models based on Voronoi cells (e.g. Hietel et al., 2000; Serrano et al., 2005; Heß & Springel, 2010), or mesh-free finite volume approaches (Junk, 2002). The former are similar in spirit to SPH and typically maintain a constant mass per particle, whereas our scheme is really closely related to ordinary mesh codes – except that the mesh is fully dynamic.

With illustrative test problems we shall later show that the resulting formulation of hydrodynamics performs rather well on a number of test problems, featuring very high accuracy in the treatment of shocks, shear waves, and fluid instabilities. In particular, it can give better results than fixed-mesh Eulerian hydrodynamics, thanks to much reduced advection errors. It also offers much higher accuracy than SPH when an equal number of particles/cells is used, making it highly attractive as a possible alternative to currently employed SPH and AMR schemes in astrophysics.

This article is structured as follows. In Section 2, we formulate continuum hydrodynamics on the Voronoi mesh, based on a finite-volume ansatz and

a second-order accurate extension of Godunov’s method. In Section 3, we briefly discuss time integration and implementation aspects. We then turn to a discussion of a number of basic hydrodynamical tests in Section 4, chosen to highlight some of the principal advantages and properties of the new approach. We note that a more extensive discussion of code tests and of the algorithmic implementation of the new scheme in the parallel AREPO code can be found in Springel (2010a). Finally, we summarize and discuss our main findings in Section 5.

2 A finite volume discretization of the Euler equations on a moving Voronoi mesh

The Euler equations are conservation laws for mass, momentum and energy that take the form of a system of hyperbolic partial differential equation. They can be written in compact form by introducing a state vector

$$\mathbf{U} = \begin{pmatrix} \rho \\ \rho \mathbf{v} \\ \rho e \end{pmatrix} = \begin{pmatrix} \rho \\ \rho \mathbf{v} \\ \rho u + \frac{1}{2} \rho \mathbf{v}^2 \end{pmatrix} \quad (1)$$

for the fluid, where ρ is the mass density, \mathbf{v} is the velocity field, and $e = u + \mathbf{v}^2/2$ is the total energy per unit mass. u gives the thermal energy per unit mass, which for an ideal gas is fully determined by the temperature. These fluid quantities are functions of the spatial coordinates \mathbf{x} and time t , i.e. $\mathbf{U} = \mathbf{U}(\mathbf{x}, t)$, but for simplicity we will typically refrain from explicitly stating this dependence in our notation. Based on \mathbf{U} , we can define a flux function

$$\mathbf{F}(\mathbf{U}) = \begin{pmatrix} \rho \mathbf{v} \\ \rho \mathbf{v} \mathbf{v}^T + P \\ (\rho e + P) \mathbf{v} \end{pmatrix}, \quad (2)$$

with an equation of state

$$P = (\gamma - 1) \rho u \quad (3)$$

that gives the pressure of the fluid. The Euler equations can then be written in the compact form

$$\frac{\partial \mathbf{U}}{\partial t} + \nabla \cdot \mathbf{F} = 0, \quad (4)$$

which emphasizes their character as conservation laws for mass, momentum and energy.

Over the past decades, a large variety of different numerical approaches to solve this coupled set of partial differential equations have been developed (see Toro, 1997; LeVeque, 2002, for comprehensive expositions). We will here employ a so-called *finite-volume* strategy, in which the discretization is carried out in terms of a subdivision of the system’s volume into a finite number of

disjoint cells. The fluid's state is described by the cell-averages of the conserved quantities for these cells. In particular, integrating the fluid over the volume V_i of cell i , we can define the total mass m_i , momentum p_i and energy E_i contained in the cell as follows,

$$\mathbf{Q}_i = \begin{pmatrix} m_i \\ \mathbf{p}_i \\ E_i \end{pmatrix} = \int_{V_i} \mathbf{U} \, dV. \quad (5)$$

With the help of the Euler equations, we can calculate the rate of change of \mathbf{Q}_i in time. Converting the volume integral over the flux divergence into a surface integral over the cell results in

$$\frac{d\mathbf{Q}_i}{dt} = - \int_{\partial V_i} [\mathbf{F}(\mathbf{U}) - \mathbf{U}\mathbf{w}^T] \, d\mathbf{n}. \quad (6)$$

Here \mathbf{n} is an outward normal vector of the cell surface, and \mathbf{w} is the velocity with which each point of the boundary of the cell moves. In Eulerian codes, the mesh is taken to be static, so that $\mathbf{w} = 0$, while in a fully Lagrangian approach, the surface would move at every point with the local flow velocity, i.e. $\mathbf{w} = \mathbf{v}$. In this case, the right hand side of equation (6) formally simplifies, because then the first component of \mathbf{Q}_i , the mass, stays fixed for each cell. Unfortunately, it is normally not possible to follow the distortions of the shapes of fluid volumes exactly in multi-dimensional flows for a reasonably long time, or in other words, one cannot guarantee the condition $\mathbf{w} = \mathbf{v}$ over the entire surface. In this case, one needs to use the general formula of equation (6), as we will do in this work.

The cells of our finite volume discretization are polyhedra with flat polygonal faces (or lines in 2D). Let \mathbf{A}_{ij} describe the oriented area of the face between cells i and j (pointing from i to j). Then we can define the averaged flux across the face i - j as

$$\mathbf{F}_{ij} = \frac{1}{A_{ij}} \int_{A_{ij}} [\mathbf{F}(\mathbf{U}) - \mathbf{U}\mathbf{w}^T] \, d\mathbf{A}_{ij}, \quad (7)$$

and the Euler equations in finite-volume form become

$$\frac{d\mathbf{Q}_i}{dt} = - \sum_j A_{ij} \mathbf{F}_{ij}. \quad (8)$$

We obtain a manifestly conservative time discretization of this equation by writing it as

$$\mathbf{Q}_i^{(n+1)} = \mathbf{Q}_i^{(n)} - \Delta t \sum_j A_{ij} \hat{\mathbf{F}}_{ij}^{(n+1/2)}, \quad (9)$$

where the $\hat{\mathbf{F}}_{ij}$ are now an appropriately time-averaged approximation to the true flux \mathbf{F}_{ij} across the cell face. The notation $\mathbf{Q}_i^{(n)}$ is meant to describe the




Fig. 2. Sketch of a Voronoi mesh and the relevant geometric quantities that enter the flux calculation across a face. In (a), we show the mesh-generating points \mathbf{r}_i and \mathbf{r}_j of two cells i and j . The face between these two cells has a center-of-mass vector \mathbf{f}_{ij} , which in general will be offset from the mid-point m_{ij} of the two points. In (b), we illustrate the two velocity vectors \mathbf{w}_i and \mathbf{w}_j associated with the mesh-generating points. These are normally chosen equal to the gas velocity in the cells, but other choices are allowed too. The motion of the mesh-generating points uniquely determines the motion of the face between the cells. Only the normal velocity \mathbf{w} is however needed for the flux computation in the rotated frame x', y' .

state of the system at step n . Note that $\hat{\mathbf{F}}_{ij} = -\hat{\mathbf{F}}_{ji}$, i.e. the discretization is manifestly conservative.

Evidently, a crucial step lies in obtaining a numerical estimate of the fluxes $\hat{\mathbf{F}}_{ij}$, and a good fraction of the literature on computational fluid dynamics is concerned with this problem. This issue is particularly important since the most straightforward (and perhaps naive) approach for estimating the fluxes, namely simply approximating them as the average of the left and right cell-centered fluxes catastrophically fails and invariably leads to severe numerical integration instabilities that render such a scheme completely useless in practice.

Many modern schemes for estimating the fluxes in a stable fashion are descendants of Godunov’s method, which revolutionized the field. By solving an exact or approximate Riemann problem at cell boundaries, Godunov’s method allows the correct identification of the eigenstructure of the local solution and of the upwind direction, which is crucial for numerical stability. While Godunov’s original method offers only first order accuracy and is relatively diffusive, it can be extended to higher-order accuracy relatively simply, and in many different ways. We employ the MUSCL-Hancock scheme (van Leer, 1984; Toro, 1997; van Leer, 2006), which is a well-known and relatively simple approach for obtaining second-order accuracy in space and time. This scheme is also popular in astronomy and used in several state-of-the art Eulerian codes (e.g. Fromang et al., 2006; Mignone et al., 2007; Cunningham et al., 2009). In its basic form, the MUSCL-Hancock scheme involves a slope-limited piece-wise linear reconstruction step within each cell, a first order prediction step for the evolution over half a timestep, and finally a Riemann solver to

estimate the time-averaged inter-cell fluxes for the timestep. After the fluxes have been applied to each cell, a new averaged state of the cells is constructed. This sequence of steps in a timestep hence follows the general REA approach.

Figure 2 gives a sketch of the geometry involved in estimating the flux across the face between two Voronoi cells. Truly multidimensional Riemann solvers have been developed recently (Wendroff, 1999; Brio et al., 2001; Balsara, 2010), but it is unclear whether they can be readily adapted to our complicated face geometry. We therefore follow the common approach and calculate the flux for each face separately, treating it as an effectively one-dimensional problem. Since we do not work with Cartesian meshes, we cannot use operating splitting to deal with the individual spatial dimensions. Rather we use a method where all the fluxes are computed in one step, and are then collectively applied to calculate the change of the conserved quantities in a cell. This *unsplit* approach implicitly accounts for “corner fluxes” (Colella, 1990) needed to recover second-order accuracy through the half-step prediction step in our MUSCL-Hancock scheme (which exploits the primitive form of the Euler equations). For defining the Riemann problem normal to a cell face, we rotate the fluid state into a suitable coordinate system with the x' -axis normal to the cell face (see sketch). This defines the left and right states across the face, which we pass to an exact Riemann solver. The latter is implemented following Toro (1997) with an extension to treat vacuum states, but it could easily be substituted with an approximate Riemann solver for higher performance, if desired. We note that in multi dimensions the transverse velocities are also required in the Riemann problem in order to identify the correct upwind transverse velocity, which is important for an accurate treatment of shear. Once the flux has been calculated with the Riemann solver, we transform it back to the lab frame.

A further important point concerns the treatment of the allowed motion of cell surfaces in our scheme. In order to obtain stable upwind behavior, the Riemann problem needs to be solved *in the frame of the moving face*. This is important as the one-dimensional Riemann problem is not Galilean-invariant in the following sense: Suppose left and right state at an interface are described by (ρ_L, P_L, v_L) and (ρ_R, P_R, v_R) , for which the Riemann solver returns an interface state (ρ_F, P_F, v_F) that is the basis for the flux estimate. For example, the mass flux across the interface is then given by $\rho_F v_F$. Consider now a velocity boost v applied both to the left and the right side. The new Riemann problem is given by $(\rho_L, P_L, v_L + v)$ and $(\rho_R, P_R, v_R + v)$, and will return a flux estimate $\rho'_F v'_F$. However, in general this will yield $\rho'_F v'_F \neq \rho_F (v_F + v)$, which implies that the calculated flux vector is not Galilean invariant.

In our new hydrodynamical scheme, each timestep involves the following basic steps:

1. Calculate a new Voronoi tessellation based on the current coordinates \mathbf{r}_i of the mesh generating points. This also gives the centers-of-mass \mathbf{s}_i of

- each cell, their volumes V_i , as well as the areas A_{ij} and centers \mathbf{f}_{ij} of all faces between cells.
2. Based on the vector of conserved fluid variables \mathbf{Q}_i associated with each cell, calculate the ‘primitive’ fluid variables $\mathbf{W}_i = (\rho_i, \mathbf{v}_i, P_i)$ for each cell.
 3. Estimate the gradients of the density, of each of the velocity components, and of the pressure in each cell, and apply a slope-limiting procedure to avoid overshoots and the introduction of new extrema.
 4. Assign velocities \mathbf{w}_i to the mesh generating points.
 5. Evaluate the Courant criterion and determine a suitable timestep size Δt .
 6. For each Voronoi face, compute the flux $\hat{\mathbf{F}}_{ij}$ across it by first determining the left and right states at the midpoint of the face by linear extrapolation from the cell midpoints, and by predicting these states forward in time by half a timestep. Solve the Riemann problem in a rotated frame that is moving with the speed of the face, and transform the result back into the lab-frame.
 7. For each cell, update its conserved quantities with the total flux over its surface multiplied by the timestep, using equation (9). This yields the new state vectors $\mathbf{Q}_i^{(n+1)}$ of the conserved variables at the end of the timestep.
 8. Move the mesh-generating points with their assigned velocities for this timestep.

For the sake of definiteness, we now briefly describe the most important details of these different steps.

2.1 Gradient estimation and linear reconstruction

According to the Green-Gauss theorem, the surface integral of a scalar function over a closed volume is equal to its gradient integrated over the same volume, i.e.

$$\int_{\partial V} \phi \, d\mathbf{n} = \int_V \nabla \phi \, dV. \quad (10)$$

This suggests one possible way to estimate the mean gradient in a Voronoi cell, in the form

$$\langle \nabla \phi \rangle_i \simeq -\frac{1}{V_i} \sum_j \phi(\mathbf{f}_{ij}) \mathbf{A}_{ij}, \quad (11)$$

where $\phi(\mathbf{f}_{ij})$ is the value of ϕ at the centroid \mathbf{f}_{ij} of the face shared by cells i and j , and \mathbf{A}_{ij} is a vector normal to the face (from j to i), with length equal to the face’s area. Based on the further approximation

$$\phi(\mathbf{f}_{ij}) \simeq \frac{1}{2}(\phi_i + \phi_j), \quad (12)$$

this provides an estimate for the local gradient. Note that with the use of equation (12), the gradient of cell i only depends on the values ϕ_j of neighboring cells, but not on ϕ_i itself. While the estimate (11) can be quite generally

applied to arbitrary tessellations, due to the use of only one Gauss point per face it is also relatively inaccurate and is not exact to linear order in general.

For the special case of Voronoi cells, it is however possible to obtain a considerably better gradient estimate with little additional effort. The key is to carry out the surface integral more accurately. It can be shown (Serrano & Espa ol, 2001; Springel, 2010a) that the gradient estimate

$$\langle \nabla \phi \rangle_i = \frac{1}{V_i} \sum_{j \neq i} A_{ij} \left([\phi_j - \phi_i] \frac{\mathbf{c}_{ij}}{r_{ij}} - \frac{\phi_i + \phi_j}{2} \frac{\mathbf{r}_{ij}}{r_{ij}} \right) \quad (13)$$

is exact to linear order, independent of the locations of the mesh-generating points of the Voronoi tessellation. Here we followed the notation of Serrano & Espa ol (2001) and defined \mathbf{c}_{ij} as the vector from the midpoint between i and j to the center-of-mass of the face between i and j . Without the term involving \mathbf{c}_{ij} this gradient estimate is the same as the simpler Green-Gauss estimate. However, retaining this extra term leads to significantly better accuracy, because the gradient estimate becomes exact to linear order for arbitrary Voronoi meshes. In practice, we shall therefore always use this gradient estimation in our MUSCL-Hancock scheme for the Euler equations, where we calculate in this way gradients for the 5 primitive variables (ρ, v_x, v_y, v_z, P) that characterize each cell.

The result (13) has also an interesting relation to the formulae obtained by Serrano & Espa ol (2001) for the partial derivatives of the volume of a Voronoi cell with respect to the location of one of the points. As Serrano & Espa ol (2001) have shown, the derivative of the volume of a Voronoi cell due to the motion of a surrounding point is given by

$$\frac{\partial V_i}{\partial \mathbf{r}_j} = -A_{ij} \left(\frac{\mathbf{c}_{ij}}{r_{ij}} + \frac{\mathbf{r}_{ij}}{2r_{ij}} \right) \quad \text{for } i \neq j. \quad (14)$$

Furthermore, they show that

$$\frac{\partial V_i}{\partial \mathbf{r}_i} = - \sum_{j \neq i} \frac{\partial V_j}{\partial \mathbf{r}_i}. \quad (15)$$

Using these relations, and noting that according to the Gauss theorem we have

$$\frac{\phi_i}{V_i} \sum_{j \neq i} A_{ij} \frac{\mathbf{r}_{ij}}{r_{ij}} = 0, \quad (16)$$

because the summation is just the surface integral of a constant function, we can also write the estimate for the gradient of ϕ at \mathbf{r}_i more compactly as

$$\langle \nabla \phi \rangle_i = - \frac{1}{V_i} \sum_j \frac{\partial V_j}{\partial \mathbf{r}_i} \phi_j. \quad (17)$$

An interesting corollary of the above is that provided $\phi(\mathbf{r})$ varies only linearly, the sum

$$S = \sum_i \phi(\mathbf{r}_i) V_i \quad (18)$$

approximates the integral $\int \phi(\mathbf{r}) dV$ *exactly*, independent of the positions of the points that generate the Voronoi tessellation.

In our approach, we use the gradients estimated with equation (13) for a linear reconstruction in each cell around the center-of-mass. For example, the density at any point $\mathbf{r} \in V_i$ of a cell is estimated as

$$\rho(\mathbf{r}) = \rho_i + \langle \nabla \rho \rangle_i \cdot (\mathbf{r} - \mathbf{s}_i), \quad (19)$$

where \mathbf{s}_i is the center of mass of the cell. Note that independent of the magnitude of the gradient and the geometry of the Voronoi cell, this linear reconstruction is conservative, i.e. the total mass in the cell m_i is identical to the volume integral over the reconstruction, $m_i = \int_{V_i} \rho(\mathbf{r}) d^3r$. An alternative choice for the reference point is to choose the mesh-generating point \mathbf{r}_i instead of \mathbf{s}_i . This is the more natural choice if the cell values are known to sample the values of the underlying field at the location of the mesh-generating points, then the reconstruction is exact to linear order. However, our input quantities are cell-averages, which correspond to linear order to the values of the underlying field sampled at the center-of-masses of the cells. For this reason we prefer the center-of-mass of a cell as reference point for the reconstruction.

Nevertheless, this highlights that large spatial offsets between the center-of-mass of a cell and its mesh-generating point are a source of errors in the linear reconstruction. It is therefore desirable to use “regular” meshes if possible, where the mesh-generating points lie close to the center-of-mass; such meshes minimize the errors in the gradient estimation and the linear reconstruction. Or in other words, we would like our Voronoi meshes to be relatively close to so-called *centroidal Voronoi meshes*, where the mesh-generating points lie exactly in the center of mass of each cell. As we discuss in bit more detail later, we have developed an efficient method for steering the mesh motion such that this regularity condition can be approximately maintained at all times.

2.2 Slope limiting procedure

In smooth parts of the flow, the above reconstruction is second-order accurate. However, in order to avoid numerical instabilities the order of the reconstruction must be reduced near fluid discontinuities, such that the introduction of new extrema by over- or undershoots in the extrapolation is avoided. This is generally achieved by applying slope limiters that reduce the size of the gradients near local extrema, or by flux limiters that replace the high-order flux with a lower order version if there are steep gradients in the upstream region of the flow.

We here generalize the original MUSCL approach to an unstructured grid by enforcing monotonicity with a slope limiting of the gradients. To this end

we require that the linearly reconstructed quantities on face centroids do not exceed the maxima or minima among all neighboring cells (Barth & Jespersen, 1989). Mathematically, we replace the gradient with a slope-limited gradient

$$\langle \nabla \phi \rangle_i' = \alpha_i \langle \nabla \phi \rangle_i, \quad (20)$$

where the slope limiter $0 \leq \alpha_i \leq 1$ for each cell is computed as

$$\alpha_i = \min(1, \psi_{ij}). \quad (21)$$

Here the minimum is taken with respect to all cells j that are neighbors of cell i , and the quantity ψ_{ij} is defined as

$$\psi_{ij} = \begin{cases} (\phi_i^{\max} - \phi_i) / \Delta\phi_{ij} & \text{for } \Delta\phi_{ij} > 0 \\ (\phi_i^{\min} - \phi_i) / \Delta\phi_{ij} & \text{for } \Delta\phi_{ij} < 0 \\ 1 & \text{for } \Delta\phi_{ij} = 0 \end{cases} \quad (22)$$

where $\Delta\phi_{ij} = \langle \nabla \phi \rangle_i \cdot (\mathbf{f}_{ij} - \mathbf{s}_i)$ is the estimated change between the centroid \mathbf{f}_{ij} and the center of cell i , and $\phi_i^{\max} = \max(\phi_j)$ and $\phi_i^{\min} = \min(\phi_j)$ are the maximum and minimum values occurring for ϕ among all neighboring cells of cell i , including i itself. We note that this slope limiting scheme does not strictly enforce the total variation diminishing property, which means that (usually reasonably small) post-shock oscillations can sometimes still occur.

2.3 Setting the velocities of the mesh generators

A particular strength of the scheme we propose here is that it can be used both as an Eulerian code, and as a Lagrangian scheme. The difference lies only in the motion of the mesh-generation points. If the mesh-generating points are arranged on a Cartesian mesh and zero velocities are adopted for them, our method is identical to a second-order accurate Eulerian code on a structured grid. Of course, one can equally well choose a different layout of the points, in which case we effectively obtain an Eulerian code on an unstructured mesh. The real advantage of the new code can be realized when we allow the mesh to move, with a velocity that is tied to the local fluid speed. In this case, we obtain a Lagrangian hydrodynamical code, which has some unique and important advantages relative to an Eulerian treatment. It is however also possible to prescribe the mesh motion through an external flow field, for example in order to smoothly concentrate resolution towards particular regions of a mesh, or to realize rotating meshes. Unlike other arbitrary Lagrangian-Eulerian (ALE) fluid dynamical methods, the method proposed here does however not rely on remapping techniques to recover from distortions of the mesh once they become severe, simply because the Voronoi tessellation produced by the continuous motion of the mesh-generating points yields a mesh geometry and topology that itself changes continuously in time, without mesh-tangling effects.

The most simple and basic approach for specifying the motion of the mesh generators is to use

$$\mathbf{w}_i = \mathbf{v}_i, \quad (23)$$

i.e. the points are moved with the fluid speed of their cell. This Lagrangian ansatz is clearly appropriate for pure advection and in smooth parts of the flow. However, in this scheme there is no mechanism built in that tries to improve the regularity of the Voronoi mesh in case the mean mass per cell should develop substantial scatter around a desired mean value, or if cells with high aspect ratios occur. If desired, such tendencies of a growing mesh irregularity can be counteracted by adding corrective velocity components to the primary mesh velocities \mathbf{w}_i of equation (23). There are many different possibilities for how exactly to do this, and we consider this freedom a strength of the formalism. In Section 3.2, we will discuss a simple regularization method that we have found to be very effective.

2.4 Flux computation

An important aspect of our approach is that the specified velocities of the mesh-generating points fully determine the motion of the whole Voronoi mesh, including, in particular, the velocities of the centroids of cell faces (see sketch in Fig. 2). This allows us to calculate the Riemann problem in the rest-frame of each of the faces.

Consider one of the faces in the tessellation and call the fluid states in the two adjacent cells the ‘left’ and ‘right’ states. We first need to determine the velocity \mathbf{w} of the face based on the velocities \mathbf{w}_L and \mathbf{w}_R of the two mesh-generating points associated with the face (they are connected by a Delaunay edge). It is clear that \mathbf{w} has a primary contribution from the mean velocity $(\mathbf{w}_L + \mathbf{w}_R)/2$ of the points, but there is also a secondary contribution \mathbf{w}' from the residual motion of the two points relative to their center of mass. This residual motion is given by $\mathbf{w}'_R = -\mathbf{w}'_L = (\mathbf{w}_R - \mathbf{w}_L)/2$, and we need to determine its impact on the motion of the face centroid. The components of \mathbf{w}'_R and \mathbf{w}'_L parallel to the line connecting the centroid \mathbf{f} of the face with the midpoint \mathbf{m} of the two mesh-generating points \mathbf{r}_L and \mathbf{r}_R induce a rotation of the face around the point \mathbf{m} . We are only interested in the normal velocity component of this motion at the centroid of the face. This can be easily computed as

$$\mathbf{w}' = \frac{(\mathbf{w}_L - \mathbf{w}_R) \cdot [\mathbf{f} - (\mathbf{r}_R + \mathbf{r}_L)/2]}{|\mathbf{r}_R - \mathbf{r}_L|} \frac{(\mathbf{r}_R - \mathbf{r}_L)}{|\mathbf{r}_R - \mathbf{r}_L|}. \quad (24)$$

The full velocity \mathbf{w} of the face is then given by

$$\mathbf{w} = \frac{\mathbf{w}_R + \mathbf{w}_L}{2} + \mathbf{w}'. \quad (25)$$

We now calculate the flux across the face using the MUSCL-Hancock approach, with the important difference that we shall carry out the calculation

in the rest-frame of the face. It is convenient to do this in the primitive variables (ρ, \mathbf{v}, P) , where we first transform the lab-frame velocities of the two cells to the rest-frame of the face by subtracting \mathbf{w} ,

$$\mathbf{W}'_{L,R} = \mathbf{W}_{L,R} - \begin{pmatrix} 0 \\ \mathbf{w} \\ 0 \end{pmatrix}. \quad (26)$$

We then linearly predict the states on both side to the centroid of the face, and also predict them forward in time by half a timestep. This produces the states

$$\mathbf{W}''_{L,R} = \mathbf{W}'_{L,R} + \left. \frac{\partial \mathbf{W}'}{\partial \mathbf{r}} \right|_{L,R} (\mathbf{f} - \mathbf{s}_{L,R}) + \left. \frac{\partial \mathbf{W}'}{\partial t} \right|_{L,R} \frac{\Delta t}{2}. \quad (27)$$

The spatial derivatives $\partial \mathbf{W}' / \partial \mathbf{r}$ are known, and given by the (slope-limited) gradients of the primitive variables that are estimated as described in Section 2.1. Note that the gradients are unaffected by the change of rest-frame described by Eqn. (26). The partial time derivate $\partial \mathbf{W} / \partial t$ can be replaced by spatial derivatives as well, based on the Euler equations in primitive variables, which are given by

$$\frac{\partial \mathbf{W}}{\partial t} + \mathbf{A}(\mathbf{W}) \frac{\partial \mathbf{W}}{\partial \mathbf{r}} = 0, \quad (28)$$

where \mathbf{A} is the matrix

$$\mathbf{A}(\mathbf{W}) = \begin{pmatrix} \mathbf{v} & \rho & 0 \\ 0 & \mathbf{v} & 1/\rho \\ 0 & \gamma P & \mathbf{v} \end{pmatrix}. \quad (29)$$

Having finally obtained the states left and right of the interface, we need to turn them into a coordinate system aligned with the face, such that we can solve an effectively one-dimensional Riemann problem. The required rotation matrix $\mathbf{\Lambda}$ for the states only affects the velocity components, viz.

$$\mathbf{W}'''_{L,R} = \mathbf{\Lambda} \mathbf{W}''_{L,R} = \begin{pmatrix} 1 & 0 & 0 \\ 0 & \mathbf{\Lambda}_{3D} & 0 \\ 0 & 0 & 1 \end{pmatrix} \mathbf{W}''_{L,R}, \quad (30)$$

where $\mathbf{\Lambda}_{3D}$ is an ordinary rotation of the coordinate system, such that the new x -axis is parallel to the normal vector of the face, pointing from the left to the right state.

With these final states, we now solve the Riemann problem, and sample the self-similar solution along $x/t = 0$. This can be written as

$$\mathbf{W}_F = R_{\text{iemann}}(\mathbf{W}'''_L, \mathbf{W}'''_R), \quad (31)$$

where R_{iemann} is a one-dimensional Riemann solver, which returns a solution for the state of the fluid \mathbf{W}_F on the face in primitive variables. We now transform this back to the lab-frame, reversing the steps above,

$$\mathbf{W}_{\text{lab}} = \begin{pmatrix} \rho \\ \mathbf{v}_{\text{lab}} \\ P \end{pmatrix} = \Lambda^{-1} \mathbf{W}_{\text{F}} + \begin{pmatrix} 0 \\ \mathbf{w} \\ 0 \end{pmatrix}. \quad (32)$$

Finally, we can use this state to calculate the fluxes in the conserved variables across the face. Here we need to take into account that the face is moving with velocity \mathbf{w} , meaning that the appropriate flux vector in the lab frame is given by

$$\hat{\mathbf{F}} = \mathbf{F}(\mathbf{U}) - \mathbf{U} \mathbf{w}^T = \begin{pmatrix} \rho(\mathbf{v}_{\text{lab}} - \mathbf{w}) \\ \rho \mathbf{v}_{\text{lab}}(\mathbf{v}_{\text{lab}} - \mathbf{w})^T + P \\ \rho e_{\text{lab}}(\mathbf{v}_{\text{lab}} - \mathbf{w}) + P \mathbf{v}_{\text{lab}} \end{pmatrix}, \quad (33)$$

where \mathbf{U} is the state \mathbf{W}_{lab} expressed in the conserved variables, and $e_{\text{lab}} = \mathbf{v}_{\text{lab}}^2/2 + P_{\text{lab}}/[(\gamma - 1)\rho_{\text{lab}}]$. The scalar product of this flux vector with the normal vector of the face gives the net flux of mass, momentum, and energy that the two adjacent, moving cells exchange. It is the flux of equation (33) that can finally be used in the conservative updates of each cell, as described by equation (9).

For the above formulation, it is straightforward to show that the changes in the conserved quantities in a cell are Galilean invariant; any Galilean boost is effectively simply absorbed into the motion of the face.

3 Time integration and implementation aspects

3.1 Time integration

For hydrodynamics with a global timestep, we employ a simplified CFL timestep criterion in the form

$$\Delta t_i = C_{\text{CFL}} \frac{R_i}{c_i + |\mathbf{v}'_i|} \quad (34)$$

to determine the maximum allowed timestep for a cell i . Here R_i is the effective radius of the cell, calculated as $R_i = (3V_i/4\pi)^{1/3}$ from the volume of a cell (or as $R_i = (V_i/\pi)^{1/2}$ from the area in 2D), under the simplifying assumption that the cell is spherical. The latter is normally a good approximation, because we steer the mesh motion such that the cell-generating point lies close to the center-of-mass of the cell, which gives it a “roundish” polyhedral shape. $C_{\text{CFL}} < 1$ is the Courant-Friedrichs-Levy coefficient (usually we choose $C_{\text{CFL}} \simeq 0.4 - 0.8$), $c_i = \sqrt{\gamma P/\rho}$ is the sound speed in the cell, and $|\mathbf{v}'_i| = |\mathbf{v}_i - \mathbf{w}_i|$ is the velocity of the gas *relative to the motion of the grid*. In the Lagrangian mode of the scheme, the velocity $|\mathbf{v}'_i|$ is close to zero and usually negligible against the sound speed, which means that larger timesteps than in an Eulerian treatment are possible, especially if there are large bulk velocities in the system.

If the code is operated with a global timestep, we determine the next system timestep as the minimum

$$\Delta t = \min_i \Delta t_i \quad (35)$$

of the timestep limits of all particles.

It is also possible to implement an individual timestep scheme, where the timestep conditions of different cells are treated in a more flexible fashion. This can greatly improve the computational efficiency in many applications. For example, in cosmological simulations, a large dynamic range in densities quickly occurs as a result of gravitational clustering. Accordingly, the local dynamical times can vary by orders of magnitude. It has therefore long become common practice to use individual timesteps for the collisionless N-body problem, a technique that has also been extended to hydrodynamical SPH simulations (e.g. Katz et al., 1996; Springel et al., 2001). We have implemented such a method also for the moving-mesh scheme, based on a discretization of the allowed timestep sizes into a power-of-two hierarchy. Unlike the approach taken in AMR simulations, where refined grid patches are typically as a whole subcycled in time by a constant factor, we impose no such restriction on our mesh, i.e. in principle each cell can be evolved with its own timestep, constrained only to the power-of-two hierarchy of allowed timestep sizes. To maintain a fully conservative character of the scheme, we evolve each face with the smaller timestep of the two adjacent cells. Full details of this individual timestep scheme can be found in Springel (2010a).

3.2 Mesh regularity

As seen in Figure 1, Voronoi meshes may sometimes look quite “irregular”, in the sense that there is a significant spread in sizes and aspect ratios of the cells, especially for disordered point distributions. While this is not a problem of principle for our approach, it is clear that the computational efficiency will normally be optimized if regions of similar gas properties are represented with cells of comparable size. Having a mixture of cells of greatly different volumes to represent a gas of constant density will restrict the size of the timestep unnecessarily (which is determined by the smallest cells), without giving any benefit in spatial resolution (which will be limited by the largest cells in the region).

As we have seen, it is also desirable to have cells where the center-of-mass lies close to the mesh-generating point, because this minimizes errors in the linear reconstruction and limits the rate at which mesh faces turn their orientation during mesh motion. Below, we will discuss one possible approach for steering the mesh motion during the dynamical evolution such that, if desired, mesh regularity in the above sense can be achieved and maintained.

In so-called centroidal Voronoi tessellations (Okabe et al., 2000), the mesh-generating points coincide with the center-of-mass of all cells. There is an

amazingly simple algorithm known as Lloyd’s method (Lloyd, 1982) to obtain a centroidal Voronoi tessellation starting from an arbitrary tessellation. One simply moves the mesh-generating points of the current Voronoi tessellation to the center-of-masses of their cells, and then reconstructs the Voronoi tessellation. The process is repeated iteratively, and with each iteration, the mesh relaxes more towards a configuration in which the Voronoi cells appear quite ‘round’ (in 2D they form a honeycomb-like mesh) and have similar volume – a centroidal Voronoi tessellation.

Inspired by this algorithm, we employ a simple scheme to improve, if needed, the local shape of the Voronoi tessellation during the dynamical evolution. We simply augment equation (23) with an additional velocity component, which is designed to move a given mesh-generating point towards the center-of-mass of its cell. There are different possibilities to parameterize such a corrective velocity. One approach that we found to work quite well in practice is to add a correction velocity whenever the mesh-generating point is further away from the center-of-mass of a cell than a given threshold, irrespective of the actual velocity field of the gas. To this end, we associate a radius $R_i = (3V_i/4\pi)^{1/3}$ with a cell based on its volume (or area in 2D). If the distance d_i between the cell’s center-of-mass \mathbf{s}_i and its mesh-generating point \mathbf{r}_i exceeds some fraction η of the cell radius R_i , we add a corrective term proportional to the local sound speed c_i of the cell to the velocity of the mesh-generating point. This effectively applies one Lloyd iteration (or a fraction of it) to the cell by repositioning the mesh-generating point onto the current center-of-mass, ignoring other components of the mesh motion. In order to soften the transition between no correction and the full correction, we parameterize the velocity as

$$\mathbf{w}'_i = \mathbf{w}_i + \chi \begin{cases} 0 & \text{for } d_i/(\eta R_i) < 0.9 \\ c_i \frac{\mathbf{s}_i - \mathbf{r}_i}{d_i} \frac{d_i - 0.9\eta R_i}{0.2\eta R_i} & \text{for } 0.9 \leq d_i/(\eta R_i) < 1.1 \\ c_i \frac{\mathbf{s}_i - \mathbf{r}_i}{d_i} & \text{for } 1.1 \leq d_i/(\eta R_i) \end{cases} \quad (36)$$

but the detailed width of this transition is unimportant. In very cold flows the sound speed may be so low that the correction becomes ineffective. As an alternative, we therefore also implemented an option in our code that allows a replacement of $c_s(\mathbf{s}_i - \mathbf{r}_i)/d_i$ in equation (36) with $(\mathbf{s}_i - \mathbf{r}_i)/\Delta t$. This more aggressive approach to ensure round cells generally works very well too, but has the disadvantage to depend on the timestepping. Our typical choice for the threshold of the correction is $\eta = 0.25$, and we usually set $\chi = 1.0$, i.e. the correction is, if present, applied in full over the course of one timestep. Smaller values of η can be used to enforce round cell shapes more aggressively, if desired.

The above scheme is usually quite effective in maintaining low aspect ratios and a regular mesh at all times during the evolution. However, the criterion for detecting cells that should get a correction velocity is not triggered if a mesh is strongly stretched or compressed in one direction; then the centers of

mass of cells can still be close to their mesh-generating points, but the aspect ratio of cells can be very high. We have found (Vogelsberger et al., 2011) that a simple alternative criterion treats such situations much better. To this end we determine for each cell the maximum angle under which any of the faces of the cell is seen from its mesh-generating point. If this angle lies above a prescribed threshold value, the mesh-correction component to the velocity is invoked, just as above. This approach will effectively try to prevent that a mesh-generating point gets too close to an outer wall of a cell, which simultaneously ensures that the displacement from the center-of-mass and the aspect ratio stay small. We have also found that with this criterion the mesh-correction motions are required more rarely, hence we have made this our default choice for general simulations with the moving mesh approach. In any case, it is important to note that the correction velocities are still Galilean-invariant, and they vanish most of the time, so that the mesh-generating points will usually be strictly advected with the local fluid velocity.

We point out that there is an important difference of this approach compared with the mesh regularization technique presented in Heß & Springel (2010). In the finite volume approach discussed here, one may in principle move the mesh-generating points in nearly arbitrary ways. Maintaining a good mesh is therefore comparatively straightforward, as described above. In contrast, the Voronoi particle model of Heß & Springel (2010) dictates a particular equation-of-motion for the mesh-generating points, where one is not allowed to simply add some mesh correction velocities. As a way out, Heß & Springel (2010) suggested to modify the underlying Lagrangian in a tricky way in order to automatically build in corrective motions into the dynamics of the mesh-generating points, but this approach is not equally flexible as the one we can use here.

3.3 Implementation aspects

The scheme described thus far has been implemented in the AREPO code, which is described in detail in Springel (2010a). A central aspect of the code is a fast engine for the generation of Delaunay and Voronoi meshes. To this end an incremental insertion algorithm is used both in 2D and 3D, which also allows partial mesh constructions, as needed in our individual timestep approach. For reasons of memory and run-time efficiency, we have written our own low-level mesh-construction routines instead of using a library such as CGAL. The mesh construction is parallelized for distributed memory machines. We use a spatial domain decomposition in which each domain is mapped to a single processor, which then first constructs its part of the mesh independently of the other CPUs, and then exchanges and inserts additional ‘ghost’ particles as needed to make sure that the Voronoi cells of all local particles are complete, i.e. that their geometry is identical to the one that would be found in a fiducial global mesh constructed in serial.

In order to robustly treat degenerate cases (for example when more than three points lie on a common circle), we work with a computational volume that is mapped to double precision numbers in the interval $[1, 2]$. For IEEE arithmetic, the mantissa of these numbers effectively defines a one-to-one mapping of all representable floating point numbers in this range to the space of 53-bit integers. We then evaluate geometric predicates with fast ordinary double precision arithmetic, but always monitor the maximum round-off error. If the outcome of a predicate is not guaranteed to be correct as a result of round-off errors, we compute the predicate exactly with long-integer arithmetic based on the mantissas corresponding to the floating point numbers. We find that this approach is both robust and still quite fast.

Finally, we would like to mention that our moving-mesh approach can quite easily be coupled to self-gravity using similar algorithms as are often employed in particle-based SPH codes. In fact, in the AREPO code we use a similar TreePM solver for gravity as employed in the GADGET code (Springel, 2005). In our approach, the Voronoi cells are treated effectively as point masses with a gravitational softening length set equal to the fiducial radius of the cell, as estimated from its volume. The tree-code has the advantage of being highly efficient also for strongly clustered particle configurations, and it can be easily adapted to individual timestep integration.

4 Illustrative test problems

We now discuss a number of simple test problems that show the performance of the moving mesh approach and illustrate its specific strengths. In some cases, we will compare directly to SPH simulations based on the same initial conditions. Also, we discuss differences in the solutions when the mesh is instead kept stationary, in which case our method behaves equivalently to a standard Eulerian scheme with second-order accuracy in space and time.

4.1 Riemann and Sod-shock problems

Arguably the most important basic test problems of hydrodynamical codes consist of one-dimensional Riemann problems. In the Riemann problem, two piece-wise constant states, each characterized by density, pressure and velocity, are brought into contact with each other, and their subsequent time evolution is then followed. If the initial velocities are zero, one deals with the special case of a Sod-shock problem. The Riemann problem can be solved analytically for an ideal gas, and the solution consists of a set of three self-similar waves that emerge from the initial discontinuity. There is in general one contact wave in the middle, sandwiched on either side by either a shock or a rarefaction. The ability of a hydrodynamical method to accurately treat different Riemann problems is fundamental for the ability of the scheme to capture complex hydrodynamical phenomena.




Fig. 3. Different one-dimensional Riemann problems, calculated with a resolution of 100 points in the unit domain, for a gas with adiabatic index $\gamma = 1.4$. The three columns show results for the initial conditions of the problems 1, 2 and 3 as specified in the text. Symbols represent the hydrodynamical quantities of the Voronoi cells, while the solid lines give the analytic solutions for density, velocity and pressure, from top to bottom.

In order to highlight the ability of our moving Voronoi mesh to accurately represent Riemann problems we consider in Fig. 3 the results for three Riemann problems, simulated in 1D with 100 initially equally spaced points in the unit domain. For problem one, the initial conditions are given by $(\rho_1, P_1, v_1, \rho_2, P_2, v_2) = (1.0, 1.0, 1.0, 0.125, 0.1, 0.0)$, for problem 2 the corresponding values are $(1.0, 0.4, -2.0, 1.0, 0.4, 2.0)$, and for problem 3 they are $(1.0, 1000.0, 0.0, 1.0, 0.01, 0.0)$. These values describe the same problems as discussed in the book by Toro (1997), and correspond to a moderate Sod shock tube, a strong double rarefaction, and a very strong shock.

As we can see in Fig. 3 from the comparison to the analytic solution, the moving mesh approach captures the solutions of these Riemann problems rather accurately. The contact discontinuities and shocks are quite sharp, with only negligible post-shock oscillations. The only significant error occurs in the nearly evacuated region of the strong rarefaction, where the simulated temperature is too high. However, this is a common error of Eulerian codes when applied to this problem. We also see that the spatial resolution varies at the end, since the points have moved with the flow. In particular, the resolution has become quite low in the low density region that develops in the middle of the double rarefaction, while it has increased on the right hand side of the contact discontinuities in the two Sod-shock problems. We note that the accuracy with which the analytic solution is recovered is considerably better than with SPH for the same initial conditions (see Springel, 2010b).

4.2 Isentropic vortex convergence

We next turn to a test of a non-trivial multi-dimensional fluid problem with a smooth solution, the isentropic vortex problem (Yee et al., 2000; Calder et al., 2002). This problem is particularly useful for verifying whether our method does indeed show second-order convergence, despite the presence of strong deformations and topological changes of the mesh and the use of small velocity components to keep the mesh nice and regular, as described above. Previously, second-order convergence of the AREPO code has only been explicitly demonstrated for 1D sound waves (Springel, 2010a), which is a comparatively simple problem where no mesh twisting occurs. The 2D Gresho vortex test on the other hand (Gresho & Chan, 1990; Liska & Wendroff, 2003) shows only a convergence rate of -1.4 as a function of the number of cells per dimension (Springel, 2010a), both in our moving mesh code and other fixed mesh codes that have second-order accuracy, like ATHENA (Stone et al., 2008). However, this can be understood as a result of the presence of discontinuities in the Gresho problem.

Yee et al. (2000) describe the setup of a perfectly smooth vortex, which has an analytic, time-invariant solution. To realize this problem, we adopt a box of extension $[-5, 5]^2$ in 2D, with periodic boundaries everywhere. The initial distribution of mesh-generating points is adopted as a regular Cartesian grid. The velocity field is specified as

$$v_x(x, y) = -y \frac{\beta}{2\pi} \exp\left(\frac{1-r^2}{2}\right) \quad (37)$$

$$v_y(x, y) = x \frac{\beta}{2\pi} \exp\left(\frac{1-r^2}{2}\right) \quad (38)$$

where $r^2 = x^2 + y^2$. The density and thermal energy per unit mass are calculated from

$$T(x, y) \equiv P/\rho = T_\infty - \frac{(\gamma-1)\beta}{8\gamma\pi^2} \exp(1-r^2) \quad (39)$$




Fig. 4. L2 error norm at time $t = 8.0$ as a function of resolution for the isentropic vortex test (left panel). The dashed line is a power law with $L2 \propto N^{-2}$. The panels on the right show the mesh in the region $[-4, 4] \times [-4, 4]$ around the center at times $t = 1.5$ (top) and $t = 8.0$ (bottom), for the run at resolution 80×80 .

as $\rho = T^{1/(\gamma-1)}$ and $u = T/(\gamma - 1)$. For these choices, the entropy P/ρ^γ is exactly constant everywhere, and the solution is time-independent. We adopt $\gamma = 1.4$, a vortex strength $\beta = 5.0$, and $T_\infty = 1$. We realize the initial conditions by integrating over the fields in each grid cell to obtain the conserved variables, as in Calder et al. (2002). We then calculate the evolution of the vortex with our moving mesh code for different resolutions until time $t = 8.0$, at which point the vortex has rotated more than once, and the mesh has been thoroughly sheared in the region of the vortex.

In Figure 4, we consider the L2-norm of the numerically obtained density field at the final time relative to the analytic solution, as a function of resolution. We use $N^2 = 40^2, 80^2, 160^2, 320^2, 640^2$, and 1280^2 initial mesh cells.

Reassuringly, the error declines accurately as a powerlaw, with $L2 \propto N^{-2}$, which is the expected convergence rate for a second-order accurate scheme. This convergence rate has also been reached by Calder et al. (2002) for the FLASH code, but unlike for this code, the error in our approach is completely independent on whether or not the vortex has an additional bulk velocity. We note that our result also disagrees with the conjecture that an additional (second) mesh-construction per time step would be needed to reach this convergence rate for multi-dimensional flow (Duffell & MacFadyen, 2011).

4.3 Rayleigh-Taylor instabilities

In multi-dimensional flows, a further class of important hydrodynamical phenomena besides acoustic waves and the non-linear waves related to Riemann problems (shocks, contact discontinuities and rarefaction waves) appears. These are so-called fluid instabilities, such as the Rayleigh-Taylor or Kelvin-Helmholtz instabilities. They are highly important for producing turbulence and for inducing mixing processes between different phases of fluids.

The Rayleigh-Taylor instability can arise in stratified layers of gas in an external gravitational field. If higher density gas lies on top of low-density gas, the stratification is unstable to buoyancy forces, and characteristic finger-like perturbations grow that will mix the fluids with time. To illustrate this instability and simultaneously show the motion of the mesh in our Voronoi based code, we illustrate in Figure 5 the evolution of a single Rayleigh-Taylor mode, calculated at the deliberately low resolution of 12×36 cells. The simulation domain is two-dimensional, with extension $[0.5, 1.5]$ and periodic boundaries at the vertical boundaries, and solid walls at the bottom and top. There is an external gravitational field with acceleration $g = -0.1$, and the bottom and top halves of the box are filled with gas of density $\rho = 1$ and $\rho = 2$, respectively. The gravitational forces are balanced by an initial hydrostatic pressure profile of the form $P(y) = P_0 + (y - 0.75) g \rho(y)$ with P_0 and $\gamma = 1.4$. To seed the perturbation, one mode is excited with a small velocity perturbation of the form $v_y(x, y) = w_0[1 - \cos(4\pi x)][1 - \cos(4\pi y/3)]$, where $w_0 = 0.0025$.

As can be clearly seen in the time evolution shown in Figure 5, the Rayleigh-Taylor instability is captured well by the moving-mesh method even at this low resolution. What is particularly interesting is that the sharp boundary between the phases can be maintained for relatively long time during the early evolution of the instability, simply because the contact discontinuity is not smeared out as it bends, thanks to the mesh's ability to follow this motion in an approximately Lagrangian fashion. A Eulerian approach with a stationary mesh on the other hand would automatically wash out the boundary due to advection errors, involving some spurious mixing of the fluids.

This fundamental improvement of the moving mesh code with respect to a fixed mesh approach becomes clearer in Figure 6. Here we compare a high-resolution version of the Rayleigh-Taylor instability between the moving-mesh approach and the same calculation carried out with a stationary Cartesian




Fig. 5. Rayleigh-Taylor instability calculated at low resolution with the moving-mesh approach. A denser fluid lies above a less dense fluid in an external gravitational field. The hydrostatic equilibrium of the initial state is buoyantly unstable. The three frames show the time evolution of the density field of the system at times $t = 5.0$, 10.0 , and 15.0 , after a single mode has been perturbed to trigger the instability, as described in the text.

mesh. Here 1024×1024 cells have been used in the unit domain, $[-0.5, 0.5]^2$, and the instability was triggered by adding small random noise to the y -velocity field, of the form $v_y(x, y) = A [1 + \cos(2\pi y)]/2$, where A is a random number in the interval $[-0.05, 0.05]$. While the instability shows a similar overall growth rate in both cases, eventually leading to full turbulence in the box, there are also striking differences. Whereas the calculation with the fixed mesh produces a lot of intermediate density values due to the strong mixing of the phases on small scales, the moving mesh approach maintains finely stratified regions where different layers of the fluid phases have been folded over each other. The contact discontinuities between these layers can be kept

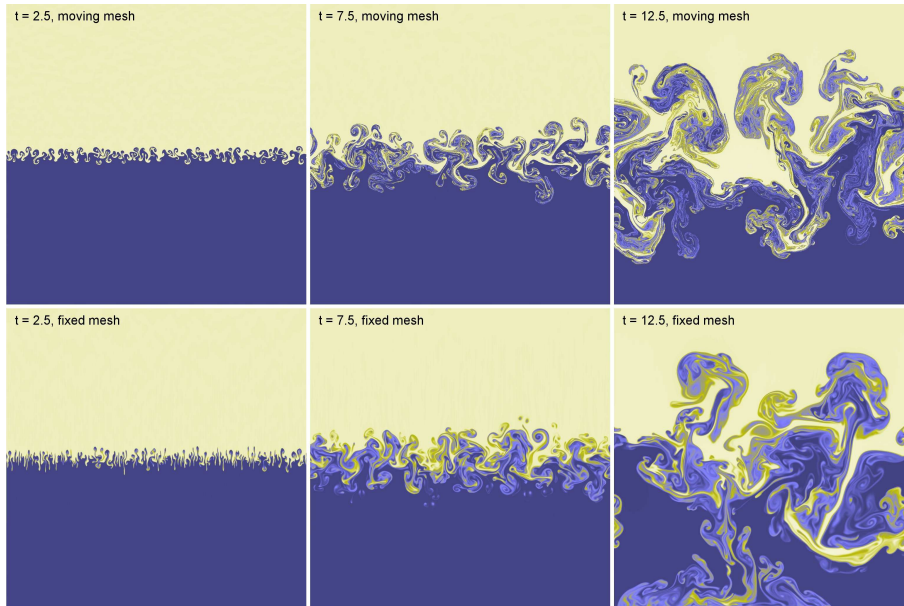


Fig. 6. Rayleigh-Taylor instability calculated at high resolution with 1024×1024 points in the unit domain. The instability is here seeded by small random noise added to the velocity field. The top and bottom rows compare the time evolution for calculations with a moving and a stationary mesh, respectively.

sharp by the code even when they are moving relative to the rest-frame of the box. We think that this behavior is much more faithful to the underlying hydrodynamical flow. In the early phase of the growth, it also appears as if small-scale RT fingers grow somewhat too quickly in the Eulerian case as a result of grid alignment effects.

4.4 Kelvin-Helmholtz instabilities

Another important fluid instability arises in shear flows, the so-called Kelvin-Helmholtz (KH) instability. Whenever there is a discontinuity in the shear velocity across a fluid interface, wave-like transverse perturbations across the interface will grow into characteristic wave-like patterns. This instability is ultimately behind the generation of waves on lakes and oceans when wind streams over the water. The KH instability is ubiquitous in complex flows and plays a prominent role in the generation of turbulence.

It has recently been found that the simulation of KH instabilities can be quite problematic in SPH, with the growth being suppressed when the density jump across the fluid discontinuity is large (Agertz et al., 2007). This has triggered a flurry of activity in the recent literature on SPH, trying to




Fig. 7. Growth rate of the KH instability for two different density ratios, 1 : 2 (left panel) or 1 : 10 (right panel). The red lines show the results for the moving mesh code when initial conditions with a sharp density jump are used, either with an initial perturbation amplitude of $v_0 = 0.001$ or $v_0 = 0.01$. The dashed blue lines are the corresponding results for a stationary mesh. The dotted lines give the exponential growth expected from linear perturbation theory. Finally, the thin red and thin blue dashed lines are the results obtained when the initial discontinuity is washed out in the initial conditions. Time is given in units of the KH growth timescale τ_{KH} .

improve on this behavior (Price, 2008; Read et al., 2010; Heß & Springel, 2010; Junk et al., 2010; Abel, 2011).

We here show a basic KH test in a two-dimensional set-up, comparing our moving-mesh approach against the traditional fixed-mesh approach. For definiteness, we fill a box with periodic boundaries and unit length on a side with gas of density $\rho_2 = 2$ in the horizontal middle stripe, and the rest with density $\rho_1 = 1$. The middle region is moving to the right with velocity $v_x = 0.5$, the rest of the gas moves to the left with velocity $v_x = -0.5$. The initial distribution of the mesh-generation points is a Cartesian grid of resolution 256×256 . We seed an initial perturbation by adding an additional component

$$v_y(x, y) = v_0 \sin(kx) \quad (40)$$

to the velocity field, where v_0 is a small number. We choose $k = 2 \times (2\pi/L)$. Hence the Fourier spectrum of the v_y field contains in the beginning only the $k_x = 2$ mode. The growth of this mode for $t > 0$ can then be conveniently measured through Fourier transforms of the velocity field.

In the left panel of Figure 7, a number of measurements of the numerical KH growth rate for a density ratio of 1 : 2 are summarized. The red lines show the result for our new moving-mesh method; the upper line is for an initial perturbation amplitude of $v_0 = 0.01$, while the lower line is for $v_0 = 0.001$. The blue dashed lines correspond in both cases to a fixed Cartesian mesh of the same resolution. The linear theory growth rate $v_y \propto \exp(t/\tau_{\text{KH}})$ is shown as dotted lines, where

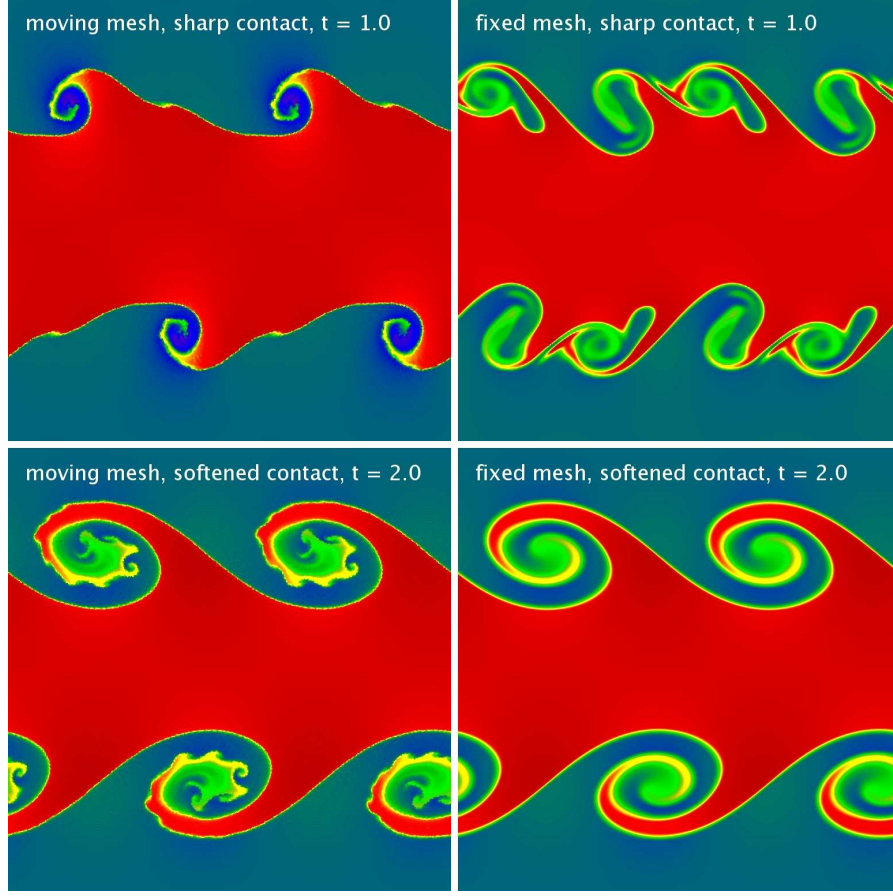


Fig. 8. Kelvin Helmholtz instability computed with different initial conditions, and for a moving or a fixed mesh. In the panels on top, the initial contact discontinuity was sharp between adjacent cells. In contrast, in the bottom row it was smoothed out.

$$\tau_{\text{KH}} = \frac{\rho_1 + \rho_2}{|v_2 - v_1| k \sqrt{\rho_1 \rho_2}} \quad (41)$$

is the KH growth timescale for an inviscid gas. Because the density is initially not perturbed self-consistently with the velocity field, it takes first a bit of time before the instability develops, but then the moving-mesh solution follows the expected linear theory growth rate quite nicely for a while. Eventually the growth slows down as the mode saturates and the non-linear evolution of the KH instability ensues.

In the fixed-mesh case, the results are somewhat less clean and depend more strongly on the initial perturbation amplitude. Visual inspection of density maps during the time evolution reveals that not only the excited mode

starts to grow but also shorter wavelength modes. This is not too surprising since small wavelength perturbations grow fastest in the Kelvin-Helmholtz instability, and representing the initial sharp density jump implicitly involves a spectrum of small waves. The latter are prone to outgrow the larger-scale seed perturbation once the discontinuity starts to be misaligned with the principal coordinate axes, in which case additional small-scale perturbations are seeded at mesh corners.

This effect can be repaired if the initial contact discontinuity is washed out, as advocated by Robertson et al. (2010). Then also in the fixed-mesh case only the excited mode grows and a more stable result is obtained. The latter is shown as thin dashed blue line in Figure 7. However, in this case one does not reach the full growth rate that is expected analytically for the (sharp) instability with this wavenumber. If one also applies the same smoothing to the initial conditions of the moving-mesh run, one obtains essentially the same result for the growth rate of the excited mode, which is shown with a thin red line. As soon as the initial discontinuity is smooth enough to be resolved by several mesh cells, it hence appears as if it would not make a difference whether one uses a moving or a fixed-mesh. This is however not true. If one waits long enough it is seen that the moving-mesh code resolves secondary KH billows for which the fixed-mesh approach appears to be already too diffusive. This can be seen in the density maps of Figure 8, where the bottom two panels compare the density field of the KH test at time $t = 2.0$ for the moving and the fixed mesh approaches, using smoothed initial conditions. The top two panels on the other hand give the same comparison when a sharp initial discontinuity is used instead. In the latter case, it is clearly seen that more ‘wrong’ modes grow in the simulation with a stationary mesh, because here a misalignment of the sharp boundary with the mesh triggers larger seed perturbations on small scales than for the moving mesh. The question whether this initial condition is somehow ‘allowed’ or not (Robertson et al., 2010) is moot in our view. Both codes are started from identical initial conditions, and hence the comparison tests how susceptible the codes are to the growth of numerically seeded small-scale perturbations in this situation.

Finally, we have also considered the same test with a density jump of $1 : 10$, which was realized by raising the density of the middle stripe to $\rho_2 = 10$. The pressure was increased by a factor of 5 to $P = 12.5$ in order to ensure that the flow stays supersonic and the Mach number of the shear flow is not much changed. The right panel in Figure 7 shows the corresponding results for the growth rate. Qualitatively, the results closely follow those obtained for the $1 : 2$ density ratio, except that the fixed-mesh results show a markedly too fast overall growth rate in this case.

4.5 Shock-cloud interaction

We finally consider two problems that involve the interaction of strong shocks with fluid instabilities, which is important in many astrophysical applications.




Fig. 9. Comparison of the time evolution of the cloud disruption test in 2D with the moving mesh code (left column) and with SPH (right column). The panels show the density field at different times, as labeled. The same set-up and the identical initial conditions as in Springel (2005) have been used. In the $t = 4.0$ frame of the moving-mesh calculation, a small rectangle marks a region that is shown enlarged at the bottom left, with the Voronoi-mesh overlaid.

First, we repeat a test presented in the GADGET code paper (Springel, 2005) as an advanced test of SPH. Here a strong shock wave of Mach number 10 strikes an initially overdense cloud with density $\rho = 5$ that is embedded at pressure equilibrium in a tenuous hot phase of density $\rho = 1$ and pressure $P = 1$ (with $\gamma = 5/3$). Figure 9 shows the time evolution of the system in 2D, comparing the moving mesh results (left row) with the SPH result obtained with the GADGET code for the identical initial conditions (right row). As the shock strikes the cloud, it is compressed and accelerated. A complicated

system of reflected and interacting shocks develops, and in the flow around the cloud, vortices are generated by the baroclinic term. With time, these vortices tend to at least partially disrupt the cloud.

In comparing the moving-mesh and the SPH results a number of interesting observations can be made. First, the density field in the smooth regions is noticeably noisy in the SPH calculation when compared with the moving-mesh approach. Also, the shock waves are not as sharp and crisp as in the Voronoi-based code, even though the global flow features are clearly very similar in both cases. Arguably the most important difference is however that the cloud is shredded much more in the moving-mesh simulation, while the SPH result shows a large degree of coherence of the cloud debris. In fact, little “droplets” of dense gas remain that are eventually advected downstream in the SPH calculation, showing no tendency to mix further with the background gas. This is presumably related to a spurious surface tension effect in SPH across contact discontinuities with large density jumps.

In the bottom left panel of the time-sequence shown in Figure 9, we have marked a small region with a black rectangle. In order to illustrate the geometry of the Voronoi mesh in this simulation, this region is shown enlarged at the bottom of Fig. 9, with the mesh overlaid. It can be seen that the higher density region in the top right is populated with smaller Voronoi cells than the lower density region at the bottom left, as a result of the Lagrangian character of the scheme.

Finally, we turn to a related simulation problem in 3D, which has become known as the ‘blob-test’. First carried out in Agertz et al. (2007), this consists of a three-dimensional overdense sphere that is put into a low-density background gas that streams supersonically with respect to the cloud. We adopt the same parameters as in Agertz et al. (2007), and use the original mesh-based initial conditions of this test as made available on the internet¹, at three different resolutions equal to $32 \times 32 \times 64$, $64 \times 64 \times 128$, and $128 \times 128 \times 256$. Similar to the two-dimensional shock-cloud interaction problem discussed above, the supersonic head wind leads to the development of a shear-flow over the surface of the cloud, which produces disrupting KH instabilities. In Agertz et al. (2007) it was found that the SPH calculations would only lead to an incomplete destruction of the cloud, while the considered Eulerian mesh-code predicted a complete destruction of the cloud after a relatively short time. The latter was measured in terms of the mass-fraction of the original cloud that was still denser than 0.64 times the density of the cloud in the initial conditions and colder than 0.9 times the temperature of the background gas.

In Figure 10, we show our results for the remaining cloud mass fraction as a function of time for three different resolutions, both for the moving-mesh code and for the equivalent calculation with a stationary mesh. We find that the moving-mesh approach appears to give converged results already at lower resolution than the Eulerian approach. It also consistently shows a

¹ They can be downloaded at <http://www.astrosim.net>




Fig. 10. The remaining mass fraction of a dense gas blob as a function of time when it is put into a supersonic head wind, corresponding to the three-dimensional ‘blob test’ of Agertz et al. (2007). The red solid, dashed, and dotted lines show our results for the moving Voronoi-mesh, with initial resolutions of $128 \times 128 \times 256$, $64 \times 64 \times 128$, and $32 \times 32 \times 64$ cells, respectively. The blue lines give the results when the mesh is kept stationary instead.

slightly higher residual cloud mass fraction than the fixed mesh calculation. Both effects can be understood as a result of the Galilean-invariance and the considerably lower advection errors of the moving mesh code. However, it is clear that both methods are qualitatively consistent and predict a complete disruption of the cloud after a timescale of $t \sim 3 \tau_{\text{KH}}$. Because SPH gives a lower mass loss at late times and does not produce a complete disruption of the cloud (Agertz et al., 2007; Heß & Springel, 2010), this reinforces the concern that gas stripping out of dense system can be systematically underestimated in SPH.

5 Discussion

We have described a novel hydrodynamical scheme on an unstructured mesh that is constructed as the Voronoi tessellation of a finite set of mesh-generating points. The points are free to move during the time evolution, allowing the mesh to seamlessly follow the flow and to change its spatial resolution fully

adaptively. Thanks to the mathematical properties of the Voronoi tessellation there are no mesh-tangling or mesh-twisting effects since the motion of the mesh-generating points induces a continuous deformation of the mesh, without producing topological artefacts.

Our approach represents a finite volume discretization of the Euler equations on an unstructured mesh with second-order accuracy both in space and time, without the need to invoke an artificial viscosity. Unlike ordinary Eulerian codes, the new method is fully Galilean-invariant, which is a very substantial advantage especially in simulations with large bulk flows. In particular, this property implies high accuracy for contact discontinuities and leads to a substantial reduction of advection errors when compared to traditional Eulerian schemes. Indeed, these advantages of the moving-mesh approach can be readily demonstrated with test problems involving fluid instabilities or moving shock waves. The Voronoi-based approach also avoids preferred spatial directions and offers flexibility in the treatment of boundary conditions. For example, curved boundaries or moving interfaces can be readily implemented.

In this paper, we have not included a discussion of the technical aspects of implementing the scheme in the new parallel cosmological code AREPO, which is described in full detail in Springel (2010a). This implementation has already been applied to timely problems in cosmological structure formation, such as galaxy formation (Vogelsberger et al., 2011) or the formation of the first stars (Greif et al., 2011a,b). Recently, important physics extensions such as radiative transfer (Petkova & Springel, 2011) and ideal magnetohydrodynamics (Pakmor et al., 2011) have been implemented in AREPO as well. The code is hence becoming an increasingly powerful alternative to more established simulation techniques in astrophysics.

This is despite the considerable computational cost that the Voronoi mesh construction entails, and despite the complicated bookkeeping code that is required for the mesh management in parallel. In 3D hydrodynamics, our Voronoi code at present is about a factor of 2 slower for the same number of resolution elements than a SPH code (if 64 smoothing neighbours are used). Compared to a Eulerian fixed mesh code, the speed difference is about a factor 3-4 (part of this difference also stems from the about twice larger average number of faces for our polyhedral cells compared with cubical cells in the Cartesian case). A further discussion of the speed difference and the scalability of the code for large cosmological applications is given in Vogelsberger et al. (2011). We note however that once self-gravity is added, the relative speed differences are much reduced, as often a sufficiently accurate calculation of gravity over a large dynamic range is more expensive than the hydrodynamics itself. Further note that our new technique reaches a given accuracy in many problems already at a lower resolution than SPH and fixed mesh codes, outweighing its higher complexity and making it hence also attractive from the point of view of computational efficiency.

References

- Abel, T. 2011, *MNRAS*, 413, 271
- Agertz, O., Moore, B., Stadel, J., Potter, D., Miniati, F., Read, J., Mayer, L., Gawryszczak, A., Kravtsov, A., Nordlund, Å., Pearce, F., Quilis, V., Rudd, D., Springel, V., Stone, J., Tasker, E., Teyssier, R., Wadsley, J., & Walder, R. 2007, *MNRAS*, 380, 963
- Balsara, D. S. 2010, *Journal of Computational Physics*, 229, 1970
- Barth, T. J. & Jespersen, D. C. 1989, *AIAA Paper*, 89-0366
- Braun, J. & Sambridge, M. 1995, *Nature*, 376, 655
- Brio, M., Zakharian, A. R., & Webb, G. M. 2001, *Journal of Computational Physics*, 167, 177
- Calder, A. C., Fryxell, B., Plewa, T., Rosner, R., Dursi, L. J., Weirs, V. G., Dupont, T., Robey, H. F., Kane, J. O., Remington, B. A., Drake, R. P., Dimonte, G., Zingale, M., Timmes, F. X., Olson, K., Ricker, P., MacNeice, P., & Tufo, H. M. 2002, *ApJS*, 143, 201
- Colella, P. 1990, *Journal of Computational Physics*, 87, 171
- Cunningham, A. J., Frank, A., Varnière, P., Mitran, S., & Jones, T. W. 2009, *ApJS*, 182, 519
- Duffell, P. C. & MacFadyen, A. I. 2011, *ArXiv e-prints*, 1104.3562
- Frenk, C. S., White, S. D. M., Bode, P., Bond, J. R., Bryan, G. L., Cen, R., Couchman, H. M. P., Evrard, A. E., Gnedin, N., Jenkins, A., Khokhlov, A. M., Klypin, A., Navarro, J. F., Norman, M. L., Ostriker, J. P., Owen, J. M., Pearce, F. R., Pen, U.-L., Steinmetz, M., Thomas, P. A., Villumsen, J. V., Wadsley, J. W., Warren, M. S., Xu, G., & Yepes, G. 1999, *ApJ*, 525, 554
- Fromang, S., Hennebelle, P., & Teyssier, R. 2006, *A&A*, 457, 371
- Gingold, R. A. & Monaghan, J. J. 1977, *MNRAS*, 181, 375
- Gnedin, N. Y. 1995, *ApJS*, 97, 231
- Greif, T. H., Springel, V., White, S. D. M., Glover, S. C. O., Clark, P. C., Smith, R. J., Klessen, R. S., & Bromm, V. 2011a, *ApJ*, 737, 75
- Greif, T. H., White, S. D. M., Klessen, R. S., & Springel, V. 2011b, *ApJ*, 736, 147
- Gresho, P. M. & Chan, S. T. 1990, *International Journal for Numerical Methods in Fluids*, 11, 621
- Hassan, O., Probert, E. J., & Morgan, K. 1998, *International Journal for Numerical Methods in Fluids*, 27, 41
- Heß, S. & Springel, V. 2010, *MNRAS*, 406, 2289
- Hietel, D., Steiner, K., & Struckmeier, J. 2000, *Mathematical Models and Methods in Applied Sciences*, 10, 1363
- Junk, M. 2002, in *Lecture Notes in Computational Science and Engineering*, Vol. 26, *Meshfree Methods for Partial Differential Equations*, ed. M. Griebel, M. A. Schweitzer, T. J. Barth, M. Griebel, D. E. Keyes, R. M. Nieminen, D. Roose, & T. Schlick (Springer Berlin Heidelberg), 223–238

- Junk, V., Walch, S., Heitsch, F., Burkert, A., Wetzstein, M., Schartmann, M., & Price, D. 2010, *MNRAS*, 407, 1933
- Katz, N., Weinberg, D. H., & Hernquist, L. 1996, *ApJS*, 105, 19
- LeVeque, R. J. 2002, *Finite volume methods for hyperbolic systems* (Cambridge University Press)
- Liska, R. & Wendroff, B. 2003, *SIAM J. Sci. Comput.*, 25, 995
- Lloyd, S. 1982, *IEEE Trans. Inform. Theory*, 28, 129137
- Lucy, L. B. 1977, *AJ*, 82, 1013
- Mavriplis, D. J. 1997, *Annual Review of Fluid Mechanics*, 29, 473
- Mignone, A., Bodo, G., Massaglia, S., Matsakos, T., Tesileanu, O., Zanni, C., & Ferrari, A. 2007, *ApJS*, 170, 228
- Mitchell, N. L., McCarthy, I. G., Bower, R. G., Theuns, T., & Crain, R. A. 2009, *MNRAS*, 395, 180
- Monaghan, J. J. 1992, *ARA&A*, 30, 543
- Okabe, A., Boots, B., Sugihara, K., & Nok Chiu, S. 2000, *Spatial Tessellations, Concepts and Applications of Voronoi Diagrams* (Chichester: John Wiley & Sons Ltd)
- Pakmor, R., Bauer, A., & Springel, V. 2011, *MNRAS*, accepted, ArXiv e-prints, 1108.1792
- Pelupessy, F. I., Schaap, W. E., & van de Weygaert, R. 2003, *A&A*, 403, 389
- Pen, U.-L. 1998, *ApJS*, 115, 19
- Petkova, M. & Springel, V. 2011, *MNRAS*, 415, 3731
- Price, D. J. 2008, *Journal of Computational Physics*, 227, 10040
- Read, J. I., Hayfield, T., & Agertz, O. 2010, *MNRAS*, 405, 1513
- Robertson, B. E., Kravtsov, A. V., Gnedin, N. Y., Abel, T., & Rudd, D. H. 2010, *MNRAS*, 401, 2463
- Sambridge, M., Braun, J., & McQueen, H. 1995, *Geophysical Journal International*, 122, 837
- Schaap, W. E. & van de Weygaert, R. 2000, *A&A*, 363, L29
- Serrano, M. & Español, P. 2001, *Phys Rev E*, 64, 046115
- Serrano, M., Español, P., & Zuniga, E. 2005, *Journal of Statistical Physics*, 121, 133
- Springel, V. 2005, *MNRAS*, 364, 1105
- . 2010a, *MNRAS*, 401, 791
- . 2010b, *ARA&A*, 48, 391
- Springel, V., Yoshida, N., & White, S. D. M. 2001, *New Astronomy*, 6, 79
- Stone, J. M., Gardiner, T. A., Teuben, P., Hawley, J. F., & Simon, J. B. 2008, *ApJS*, 178, 137
- Tasker, E. J., Brunino, R., Mitchell, N. L., Michielsen, D., Hopton, S., Pearce, F. R., Bryan, G. L., & Theuns, T. 2008, *MNRAS*, 390, 1267
- Toro, E. 1997, *Riemann solvers and numerical methods for fluid dynamics* (Springer)
- Trac, H. & Pen, U.-L. 2004, *New Astronomy*, 9, 443
- van de Weygaert, R. 1994, *A&A*, 283, 361

- van de Weygaert, R. & Schaap, W. 2009, in *Lecture Notes in Physics*, Berlin Springer Verlag, Vol. 665, *Data Analysis in Cosmology*, ed. V. J. Martínez, E. Saar, E. Martínez-González, & M.-J. Pons-Bordería, 291–413
- van Leer, B. 1984, *SIAM J. Sci. Stat. Comput.*, 5, 1
- . 2006, *Communications in Computational Physics*, 1, 192
- Vogelsberger, M., Sijacki, D., Keres, D., Springel, V., & Hernquist, L. 2011, *ArXiv e-prints*, 1109.1281
- Wadsley, J. W., Veeravalli, G., & Couchman, H. M. P. 2008, *MNRAS*, 387, 427
- Wendroff, B. 1999, *Computers & Mathematics with Applications*, 38, 175
- Whitehurst, R. 1995, *MNRAS*, 277, 655
- Xu, G. 1997, *MNRAS*, 288, 903
- Yee, H. C., Vinokur, M., & Djomehri, M. J. 2000, *Journal of Computational Physics*, 162, 33